

car price

March 12, 2022

1 Libraries

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

import warnings
warnings.filterwarnings("ignore")
```

2 Dataset

```
[2]: df = pd.read_csv('CarPrice_Assignment.csv')
df.head()
```

```
[2]:   car_ID  symboling          CarName fueltype aspiration doornumber \
0        1          3    alfa-romero giulia      gas         std         two
1        2          3    alfa-romero stelvio      gas         std         two
2        3          1  alfa-romero Quadrifoglio      gas         std         two
3        4          2          audi 100 ls      gas         std         four
4        5          2          audi 100ls      gas         std         four
```

```
   carbody drivewheel enginelocation  wheelbase  ...  enginesize  \
0  convertible      rwd          front      88.6  ...      130
1  convertible      rwd          front      88.6  ...      130
2   hatchback      rwd          front      94.5  ...      152
3        sedan      fwd          front      99.8  ...      109
4        sedan      4wd          front      99.4  ...      136
```

```
   fuelsystem  boreratio  stroke  compressionratio  horsepower  peakrpm  citympg  \
0         mpfi       3.47    2.68              9.0          111     5000      21
1         mpfi       3.47    2.68              9.0          111     5000      21
2         mpfi       2.68    3.47              9.0          154     5000      19
3         mpfi       3.19    3.40             10.0          102     5500      24
4         mpfi       3.19    3.40              8.0          115     5500      18
```

```
   highwaympg  price
```

```

0          27  13495.0
1          27  16500.0
2          26  16500.0
3          30  13950.0
4          22  17450.0

```

[5 rows x 26 columns]

```
[3]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null    int64
1   symboling             205 non-null    int64
2   CarName               205 non-null    object
3   fueltype              205 non-null    object
4   aspiration            205 non-null    object
5   doornumber            205 non-null    object
6   carbody               205 non-null    object
7   drivewheel           205 non-null    object
8   enginelocation        205 non-null    object
9   wheelbase             205 non-null    float64
10  carlength             205 non-null    float64
11  carwidth              205 non-null    float64
12  carheight             205 non-null    float64
13  curbweight            205 non-null    int64
14  enginetype            205 non-null    object
15  cylindernumber        205 non-null    object
16  enginesize            205 non-null    int64
17  fuelsystem            205 non-null    object
18  boreratio             205 non-null    float64
19  stroke                205 non-null    float64
20  compressionratio      205 non-null    float64
21  horsepower            205 non-null    int64
22  peakrpm               205 non-null    int64
23  citympg               205 non-null    int64
24  highwaympg            205 non-null    int64
25  price                 205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB

```

```
[4]: df.columns
```

```
[4]: Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',
          'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',
          'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype',
          'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',
          'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
          'price'],
          dtype='object')
```

```
[5]: cars_numeric = df.select_dtypes(include=['float64', 'int64'])
cars_numeric.head()
```

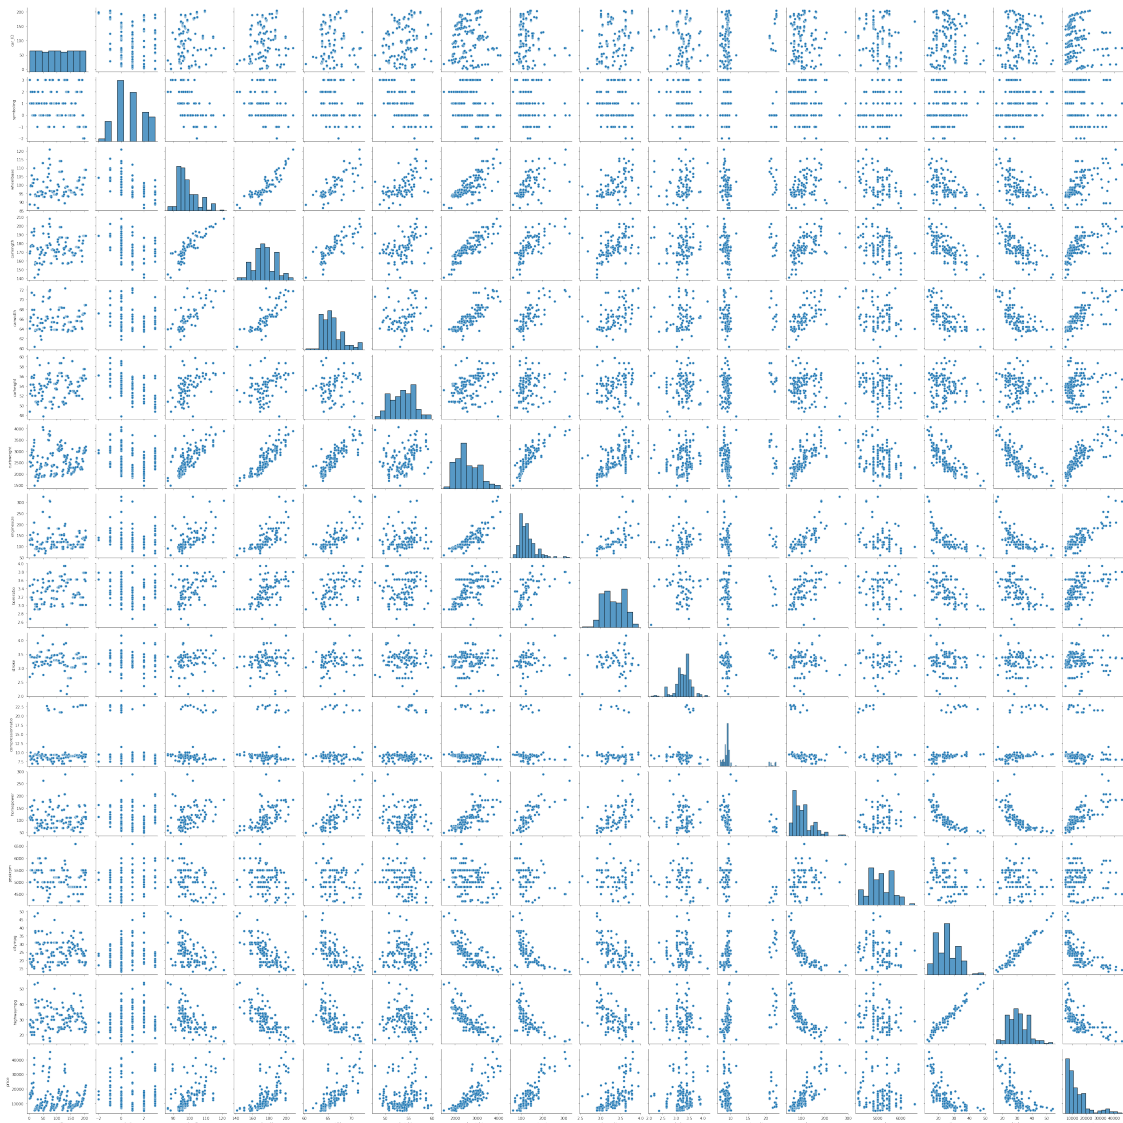
```
[5]:
```

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	\
0	1	3	88.6	168.8	64.1	48.8	2548	
1	2	3	88.6	168.8	64.1	48.8	2548	
2	3	1	94.5	171.2	65.5	52.4	2823	
3	4	2	99.8	176.6	66.2	54.3	2337	
4	5	2	99.4	176.6	66.4	54.3	2824	

	enginesize	boreratio	stroke	compressionratio	horsepower	peakrpm	\
0	130	3.47	2.68		9.0	111	5000
1	130	3.47	2.68		9.0	111	5000
2	152	2.68	3.47		9.0	154	5000
3	109	3.19	3.40		10.0	102	5500
4	136	3.19	3.40		8.0	115	5500

	citympg	highwaympg	price
0	21	27	13495.0
1	21	27	16500.0
2	19	26	16500.0
3	24	30	13950.0
4	18	22	17450.0

```
[6]: sns.pairplot(cars_numeric)
plt.show()
```



```
[7]: cars_numeric.corr()
```

```
[7]:
```

	car_ID	symboling	wheelbase	carlength	carwidth	\
car_ID	1.000000	-0.151621	0.129729	0.170636	0.052387	
symboling	-0.151621	1.000000	-0.531954	-0.357612	-0.232919	
wheelbase	0.129729	-0.531954	1.000000	0.874587	0.795144	
carlength	0.170636	-0.357612	0.874587	1.000000	0.841118	
carwidth	0.052387	-0.232919	0.795144	0.841118	1.000000	
carheight	0.255960	-0.541038	0.589435	0.491029	0.279210	
curbweight	0.071962	-0.227691	0.776386	0.877728	0.867032	
enginesize	-0.033930	-0.105790	0.569329	0.683360	0.735433	
boreratio	0.260064	-0.130051	0.488750	0.606454	0.559150	
stroke	-0.160824	-0.008735	0.160959	0.129533	0.182942	

compressionratio	0.150276	-0.178515	0.249786	0.158414	0.181129
horsepower	-0.015006	0.070873	0.353294	0.552623	0.640732
peakrpm	-0.203789	0.273606	-0.360469	-0.287242	-0.220012
citympg	0.015940	-0.035823	-0.470414	-0.670909	-0.642704
highwaympg	0.011255	0.034606	-0.544082	-0.704662	-0.677218
price	-0.109093	-0.079978	0.577816	0.682920	0.759325

	carheight	curbweight	enginesize	boreratio	stroke \
car_ID	0.255960	0.071962	-0.033930	0.260064	-0.160824
symboling	-0.541038	-0.227691	-0.105790	-0.130051	-0.008735
wheelbase	0.589435	0.776386	0.569329	0.488750	0.160959
carlength	0.491029	0.877728	0.683360	0.606454	0.129533
carwidth	0.279210	0.867032	0.735433	0.559150	0.182942
carheight	1.000000	0.295572	0.067149	0.171071	-0.055307
curbweight	0.295572	1.000000	0.850594	0.648480	0.168790
enginesize	0.067149	0.850594	1.000000	0.583774	0.203129
boreratio	0.171071	0.648480	0.583774	1.000000	-0.055909
stroke	-0.055307	0.168790	0.203129	-0.055909	1.000000
compressionratio	0.261214	0.151362	0.028971	0.005197	0.186110
horsepower	-0.108802	0.750739	0.809769	0.573677	0.080940
peakrpm	-0.320411	-0.266243	-0.244660	-0.254976	-0.067964
citympg	-0.048640	-0.757414	-0.653658	-0.584532	-0.042145
highwaympg	-0.107358	-0.797465	-0.677470	-0.587012	-0.043931
price	0.119336	0.835305	0.874145	0.553173	0.079443

	compressionratio	horsepower	peakrpm	citympg \
car_ID	0.150276	-0.015006	-0.203789	0.015940
symboling	-0.178515	0.070873	0.273606	-0.035823
wheelbase	0.249786	0.353294	-0.360469	-0.470414
carlength	0.158414	0.552623	-0.287242	-0.670909
carwidth	0.181129	0.640732	-0.220012	-0.642704
carheight	0.261214	-0.108802	-0.320411	-0.048640
curbweight	0.151362	0.750739	-0.266243	-0.757414
enginesize	0.028971	0.809769	-0.244660	-0.653658
boreratio	0.005197	0.573677	-0.254976	-0.584532
stroke	0.186110	0.080940	-0.067964	-0.042145
compressionratio	1.000000	-0.204326	-0.435741	0.324701
horsepower	-0.204326	1.000000	0.131073	-0.801456
peakrpm	-0.435741	0.131073	1.000000	-0.113544
citympg	0.324701	-0.801456	-0.113544	1.000000
highwaympg	0.265201	-0.770544	-0.054275	0.971337
price	0.067984	0.808139	-0.085267	-0.685751

	highwaympg	price
car_ID	0.011255	-0.109093
symboling	0.034606	-0.079978
wheelbase	-0.544082	0.577816

```

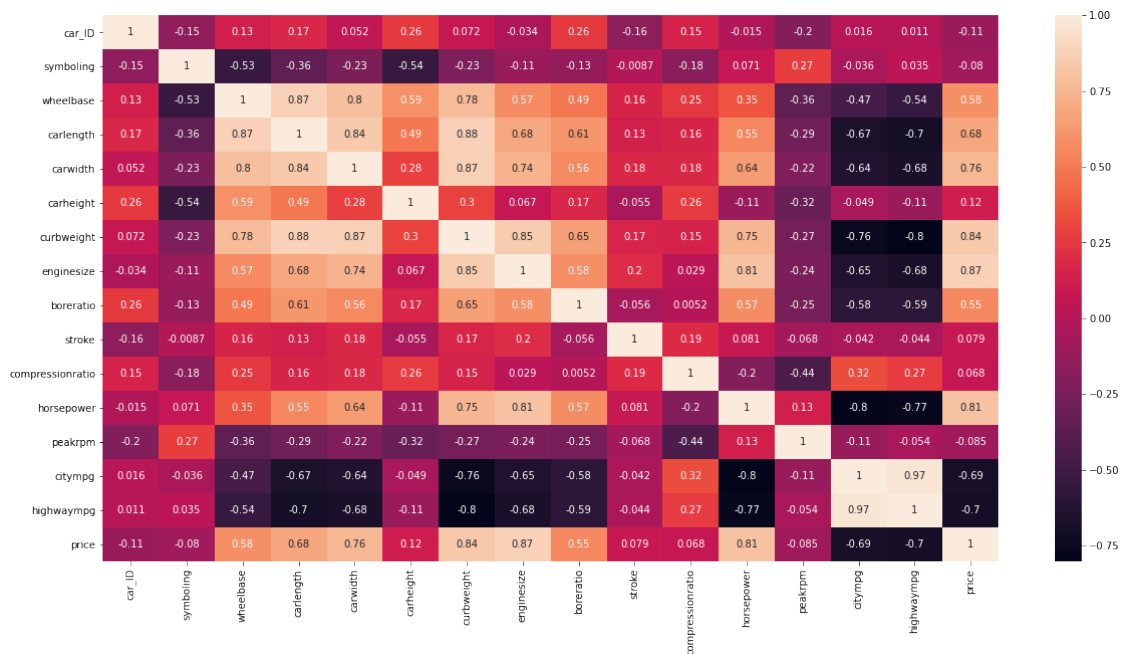
carlength      -0.704662  0.682920
carwidth       -0.677218  0.759325
carheight      -0.107358  0.119336
curbweight     -0.797465  0.835305
enginesize     -0.677470  0.874145
bore_ratio     -0.587012  0.553173
stroke         -0.043931  0.079443
compressionratio 0.265201  0.067984
horsepower     -0.770544  0.808139
peakrpm        -0.054275 -0.085267
citympg         0.971337 -0.685751
highwaympg      1.000000 -0.697599
price          -0.697599  1.000000

```

```

[8]: plt.figure(figsize=(20,10))
      sns.heatmap(cars_numeric.corr(), annot=True)
      plt.show()

```



```

[9]: df['car_company'] = df['CarName'].apply(lambda x : x.split(' ')[0])
      df.head()

```

```

[9]:   car_ID  symboling      CarName fueltype aspiration doornumber \
0      1          3    alfa-romero giulia      gas      std        two
1      2          3    alfa-romero stelvio      gas      std        two
2      3          1  alfa-romero Quadrifoglio      gas      std        two
3      4          2      audi 100 ls      gas      std        four

```

4	5	2	audi 100ls	gas	std	four
---	---	---	------------	-----	-----	------

	carbody	drivewheel	engine	location	wheelbase	...	fuelsystem	\
0	convertible	rwd	front	88.6	...		mpfi	
1	convertible	rwd	front	88.6	...		mpfi	
2	hatchback	rwd	front	94.5	...		mpfi	
3	sedan	fwd	front	99.8	...		mpfi	
4	sedan	4wd	front	99.4	...		mpfi	

	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	highwaympg	\
0	3.47	2.68	9.0	111	5000	21	27	
1	3.47	2.68	9.0	111	5000	21	27	
2	2.68	3.47	9.0	154	5000	19	26	
3	3.19	3.40	10.0	102	5500	24	30	
4	3.19	3.40	8.0	115	5500	18	22	

	price	car_company
0	13495.0	alfa-romero
1	16500.0	alfa-romero
2	16500.0	alfa-romero
3	13950.0	audi
4	17450.0	audi

[5 rows x 27 columns]

3 Spelling mistake fixing

```
[10]: df.loc[(df['car_company']=='vw') | (df['car_company']=='vokswagen'),
           ↪ 'car_company']='volkswagen'
df.loc[(df['car_company']=='toyouta'), 'car_company']='toyota'
df.loc[(df['car_company']=='maxda'), 'car_company']='mazda'
df.loc[(df['car_company']=='Nissan'), 'car_company']='nissan'
df.loc[(df['car_company']=='porcshce'), 'car_company']='porsche'
# df['car_company'] = df['car_company'].str.replace('maxda', 'mazda')
```

```
[11]: df['car_company'].value_counts()
```

```
[11]: toyota      32
      nissan      18
      mazda      17
      mitsubishi  13
      honda      13
      volkswagen  12
      subaru     12
      peugeot    11
      volvo      11
```

```
dodge          9
buick          8
bmw           8
audi          7
plymouth      7
saab          6
porsche       5
isuzu         4
jaguar        3
chevrolet     3
alfa-romero   3
renault       2
mercury       1
Name: car_company, dtype: int64
```

```
[12]: df = df.drop(['CarName'], axis=1)
```

```
[13]: df['cylindernumber'].value_counts()
```

```
[13]: four      159
      six       24
      five     11
      eight     5
      two       4
      three     1
      twelve    1
      Name: cylindernumber, dtype: int64
```

```
[14]: def number(x):
      return x.map({'four':4,
                    'six':6,
                    'five':5,
                    'eight':8,
                    'two':2,
                    'three':3,
                    'twelve':12})
```

```
[15]: df['cylindernumber'] = df[['cylindernumber']].apply(number)
```

```
[16]: df['doornumber'] = df[['doornumber']].apply(number)
```

```
[17]: df_category = df.select_dtypes(include=['object'])
      df_category
```

```
[17]:   fuelttype aspiration   carbody drivewheel enginelocation enginetype \
0      gas      std  convertible      rwd      front      dohc
1      gas      std  convertible      rwd      front      dohc
```


2	gas	std	hatchback	rwd	front	ohcv
3	gas	std	sedan	fwd	front	ohc
4	gas	std	sedan	4wd	front	ohc
..
200	gas	std	sedan	rwd	front	ohc
201	gas	turbo	sedan	rwd	front	ohc
202	gas	std	sedan	rwd	front	ohcv
203	diesel	turbo	sedan	rwd	front	ohc
204	gas	turbo	sedan	rwd	front	ohc

	fuelsystem	car_company
0	mpfi	alfa-romero
1	mpfi	alfa-romero
2	mpfi	alfa-romero
3	mpfi	audi
4	mpfi	audi
..
200	mpfi	volvo
201	mpfi	volvo
202	mpfi	volvo
203	idi	volvo
204	mpfi	volvo

[205 rows x 8 columns]

```
[18]: df_dummies = pd.get_dummies(df_category, drop_first=True)
df_dummies.head()
```

```
[18]: fueltype_gas aspiration_turbo carbody_hardtop carbody_hatchback \
0          1          0          0          0
1          1          0          0          0
2          1          0          0          1
3          1          0          0          0
4          1          0          0          0

carbody_sedan carbody_wagon drivewheel_fwd drivewheel_rwd \
0          0          0          0          1
1          0          0          0          1
2          0          0          0          1
3          1          0          1          0
4          1          0          0          0

enginelocation_rear enginetype_dohcv ... car_company_nissan \
0          0          0 ...          0
1          0          0 ...          0
2          0          0 ...          0
3          0          0 ...          0
```

```

4          0          0 ...          0

car_company_peugeot car_company_plymouth car_company_porsche \
0          0          0          0
1          0          0          0
2          0          0          0
3          0          0          0
4          0          0          0

car_company_renault car_company_saab car_company_subaru \
0          0          0          0
1          0          0          0
2          0          0          0
3          0          0          0
4          0          0          0

car_company_toyota car_company_volkswagen car_company_volvo
0          0          0          0
1          0          0          0
2          0          0          0
3          0          0          0
4          0          0          0

```

[5 rows x 43 columns]

```
[19]: df = df.drop(list(df_category.columns), axis=1)
df.head()
```

```

[19]: car_ID  symboling  doornumber  wheelbase  carlength  carwidth  carheight  \
0         1         3           2      88.6      168.8      64.1      48.8
1         2         3           2      88.6      168.8      64.1      48.8
2         3         1           2      94.5      171.2      65.5      52.4
3         4         2           4      99.8      176.6      66.2      54.3
4         5         2           4      99.4      176.6      66.4      54.3

curbweight  cylindernumber  enginesize  boreratio  stroke  \
0        2548             4         130        3.47    2.68
1        2548             4         130        3.47    2.68
2        2823             6         152        2.68    3.47
3        2337             4         109        3.19    3.40
4        2824             5         136        3.19    3.40

compressionratio  horsepower  peakrpm  citympg  highwaympg  price
0             9.0          111    5000      21          27  13495.0
1             9.0          111    5000      21          27  16500.0
2             9.0          154    5000      19          26  16500.0
3            10.0          102    5500      24          30  13950.0

```

4	8.0	115	5500	18	22	17450.0
---	-----	-----	------	----	----	---------

```
[20]: df = pd.concat([df, df_dummies], axis=1)
df.head()
```

```
[20]:
```

	car_ID	symboling	doornumber	wheelbase	carlength	carwidth	carheight	\
0	1	3	2	88.6	168.8	64.1	48.8	
1	2	3	2	88.6	168.8	64.1	48.8	
2	3	1	2	94.5	171.2	65.5	52.4	
3	4	2	4	99.8	176.6	66.2	54.3	
4	5	2	4	99.4	176.6	66.4	54.3	

	curbweight	cylindernumber	enginesize	...	car_company_nissan	\
0	2548	4	130	...	0	
1	2548	4	130	...	0	
2	2823	6	152	...	0	
3	2337	4	109	...	0	
4	2824	5	136	...	0	

	car_company_peugeot	car_company_plymouth	car_company_porsche	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	car_company_renault	car_company_saab	car_company_subaru	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	car_company_toyota	car_company_volkswagen	car_company_volvo
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

[5 rows x 61 columns]

```
[21]: 18+43
```

```
[21]: 61
```

```
[22]: df.drop(['car_ID'], axis=1, inplace=True)
```

```
[23]: df.head()
```

```
[23]:   symboling  doornumber  wheelbase  carlength  carwidth  carheight  \
0          3           2      88.6      168.8      64.1      48.8
1          3           2      88.6      168.8      64.1      48.8
2          1           2      94.5      171.2      65.5      52.4
3          2           4      99.8      176.6      66.2      54.3
4          2           4      99.4      176.6      66.4      54.3

      curbweight  cylindernumber  enginesize  boreratio  ...  car_company_nissan  \
0          2548                4          130        3.47  ...                0
1          2548                4          130        3.47  ...                0
2          2823                6          152        2.68  ...                0
3          2337                4          109        3.19  ...                0
4          2824                5          136        3.19  ...                0

      car_company_peugeot  car_company_plymouth  car_company_porsche  \
0                        0                      0                     0
1                        0                      0                     0
2                        0                      0                     0
3                        0                      0                     0
4                        0                      0                     0

      car_company_renault  car_company_saab  car_company_subaru  \
0                        0                0                    0
1                        0                0                    0
2                        0                0                    0
3                        0                0                    0
4                        0                0                    0

      car_company_toyota  car_company_volkswagen  car_company_volvo
0                        0                      0                    0
1                        0                      0                    0
2                        0                      0                    0
3                        0                      0                    0
4                        0                      0                    0

[5 rows x 60 columns]
```

4 Model Building

```
[24]: from sklearn.model_selection import train_test_split
```

```
[25]: df_train, df_test = train_test_split(df, train_size=0.7, random_state=100)
```

5 Standard Scaler

```
[26]: from sklearn.preprocessing import StandardScaler
```

```
[27]: sc = StandardScaler()
```

```
[28]: varlist = ['symboling', 'doornumber', 'wheelbase', 'carlength', 'carwidth',
                'carheight', 'curbweight', 'cylindernumber', 'enginesize',
                ↪ 'boreratio',
                'stroke', 'compressionratio', 'horsepower', 'peakrpm', 'citympg',
                'highwaympg']
```

```
[29]: df_train[varlist] = sc.fit_transform(df_train[varlist])
      #df_test = sc.transform(df_train[varlist])
```

```
[30]: df_train.head()
```

```
[30]:
```

	symboling	doornumber	wheelbase	carlength	carwidth	carheight	\
122	0.170159	0.887412	-0.811836	-0.487238	-0.924500	-1.134628	
125	1.848278	-1.126872	-0.677177	-0.359789	1.114978	-1.382026	
166	0.170159	-1.126872	-0.677177	-0.375720	-0.833856	-0.392434	
1	1.848278	-1.126872	-1.670284	-0.367754	-0.788535	-1.959288	
199	-1.507960	0.887412	0.972390	1.225364	0.616439	1.627983	

	curbweight	cylindernumber	enginesize	boreratio	...	\
122	-0.642128	-0.351431	-0.660242	-1.297329	...	
125	0.439415	-0.351431	0.637806	2.432256	...	
166	-0.441296	-0.351431	-0.660242	-0.259197	...	
1	0.015642	-0.351431	0.123485	0.625138	...	
199	1.137720	-0.351431	0.123485	1.201877	...	

	car_company_nissan	car_company_peugeot	car_company_plymouth	\
122	0	0	1	
125	0	0	0	
166	0	0	0	
1	0	0	0	
199	0	0	0	

	car_company_porsche	car_company_renault	car_company_saab	\
122	0	0	0	
125	1	0	0	
166	0	0	0	
1	0	0	0	
199	0	0	0	

	car_company_subaru	car_company_toyota	car_company_volkswagen	\
122	0	0	0	

125	0	0	0
166	0	1	0
1	0	0	0
199	0	0	0

	car_company_volvo
122	0
125	0
166	0
1	0
199	1

[5 rows x 60 columns]

```
[31]: y_train = df_train.pop('price')
```

```
[32]: X_train = df_train
```

6 Model(Liniar Regression)

```
[33]: from sklearn.linear_model import LinearRegression
```

```
[34]: lm = LinearRegression()
      lm.fit(X_train, y_train)
```

```
[34]: LinearRegression()
```

```
[35]: lm.score(X_train, y_train)
```

```
[35]: 0.9712064047413826
```

```
[36]: lm.coef_
```

```
[36]: array([-6.88567127e+01,  1.59791915e+02,  1.67620993e+03, -9.40501426e+02,
           1.66728468e+03, -1.32048230e+03,  2.11588268e+03, -2.24692316e+03,
           7.76550622e+03, -2.46121683e+03, -8.48196473e+02, -3.46772929e+03,
          -9.97487191e+02,  1.47072292e+03,  4.74859336e+02,  6.18660770e+02,
          -5.55423212e+03,  3.05402080e+03, -4.42195370e+03, -4.93031210e+03,
          -4.14820586e+03, -3.45584500e+03, -4.90769833e+02,  3.55723578e+02,
           7.66348624e+03,  7.02233939e+03,  7.73848811e+03,  2.20125928e+03,
           4.74919992e+03,  6.31832027e+01,  8.71702434e+03,  9.33222641e+02,
          -2.37609763e+03,  5.55423212e+03, -1.81898940e-12, -3.11421930e+02,
          -4.55560035e+02,  5.45696821e-12, -9.98396594e+02,  7.91156079e+03,
           8.74512513e+02, -4.59870841e+03, -5.59057700e+03, -3.90189453e+03,
          -2.38532312e+03, -1.93079078e+03, -1.22727036e+03,  0.00000000e+00,
          -6.29519219e+03, -1.93401393e+03, -1.03724527e+04, -5.55044819e+03,
           6.18313769e+03, -2.57784497e+03,  5.48694996e+03, -2.91428633e+03,
```

```
-1.38182873e+03, -1.37415291e+03, 1.45115753e+02])
```

```
[37]: lm.intercept_
```

```
[37]: 20603.961509963374
```

```
[38]: from sklearn.feature_selection import RFE
```

```
[39]: lm = LinearRegression()
rfe1 = RFE(lm, 15)
rfe1.fit(X_train, y_train)
```

```
[39]: RFE(estimator=LinearRegression(), n_features_to_select=15)
```

```
[40]: rfe1.ranking_
```

```
[40]: array([41, 39, 28, 31, 1, 26, 17, 32, 1, 18, 33, 1, 23, 24, 37, 30, 1,
        16, 14, 12, 13, 15, 35, 40, 1, 22, 1, 20, 1, 42, 1, 36, 10, 1,
        43, 27, 19, 44, 9, 1, 34, 8, 6, 4, 25, 29, 1, 45, 1, 2, 1,
        7, 11, 1, 21, 1, 3, 5, 38])
```

```
[41]: rfe1.support_
```

```
[41]: array([False, False, False, False,  True, False, False, False,  True,
        False, False,  True, False, False, False, False,  True, False,
        False, False, False, False, False, False,  True, False,  True,
        False,  True, False,  True, False, False,  True, False, False,
        False, False, False,  True, False, False, False, False, False,
        False,  True, False,  True, False,  True, False, False,  True,
        False,  True, False, False, False])
```

```
[42]: import statsmodels.api as sm
```

```
[43]: col1 = X_train.columns[rfe1.support_]
```

```
[44]: X_train_rfe1 = X_train[col1]
```

```
[45]: X_train_rfe1.head()
```

```
[45]:      carwidth  enginesize  compressionratio  fueltype_gas  \
122 -0.924500   -0.660242         -0.172569             1
125  1.114978    0.637806         -0.146125             1
166 -0.833856   -0.660242         -0.172569             1
1   -0.788535    0.123485         -0.278345             1
199  0.616439    0.123485         -0.675002             1

      enginelocation_rear  enginetype_l  enginetype_ohcf  enginetype_rotor  \
122                    0              0                0                0
```

125	0	0	0	0
166	0	0	0	0
1	0	0	0	0
199	0	0	0	0

	fuelsystem_idi	car_company_bmw	car_company_mazda	\
122	0	0	0	
125	0	0	0	
166	0	0	0	
1	0	0	0	
199	0	0	0	

	car_company_mitsubishi	car_company_peugeot	car_company_renault	\
122	0	0	0	
125	0	0	0	
166	0	0	0	
1	0	0	0	
199	0	0	0	

	car_company_subaru
122	0
125	0
166	0
1	0
199	0

```
[46]: # new columns const added
X_train_rfe1 = sm.add_constant(X_train_rfe1)
```

```
[47]: X_train_rfe1.head()
```

```
[47]:
```

	const	carwidth	enginesize	compressionratio	fueltype_gas	\
122	1.0	-0.924500	-0.660242	-0.172569	1	
125	1.0	1.114978	0.637806	-0.146125	1	
166	1.0	-0.833856	-0.660242	-0.172569	1	
1	1.0	-0.788535	0.123485	-0.278345	1	
199	1.0	0.616439	0.123485	-0.675002	1	

	engineloation_rear	enginetype_l	enginetype_ohcf	enginetype_rotor	\
122	0	0	0	0	
125	0	0	0	0	
166	0	0	0	0	
1	0	0	0	0	
199	0	0	0	0	

	fuelsystem_idi	car_company_bmw	car_company_mazda	\
122	0	0	0	

125	0	0	0
166	0	0	0
1	0	0	0
199	0	0	0

	car_company_mitsubishi	car_company_peugeot	car_company_renault \
122	0	0	0
125	0	0	0
166	0	0	0
1	0	0	0
199	0	0	0

	car_company_subaru
122	0
125	0
166	0
1	0
199	0

```
[48]: lm1 = sm.OLS(y_train, X_train_rfe1).fit()
```

```
[49]: lm1.summary()
```

```
[49]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:                  0.920
Model:                            OLS    Adj. R-squared:              0.912
Method:                 Least Squares    F-statistic:                  114.1
Date:                  Sat, 12 Mar 2022    Prob (F-statistic):          4.59e-64
Time:                  20:09:00           Log-Likelihood:              -1303.5
No. Observations:                  143    AIC:                         2635.
Df Residuals:                      129    BIC:                         2676.
Df Model:                          13
Covariance Type:                  nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const                1.242e+04    1263.541      9.829      0.000     9919.615
1.49e+04
carwidth              3348.1798     356.058      9.403      0.000     2643.710
4052.649
enginesize            3726.7427     351.763     10.594      0.000     3030.772

```

```

4422.714
compressionratio      -3504.8498   1260.300   -2.781   0.006   -5998.385
-1011.314
fueltype_gas          -525.2133   1616.959   -0.325   0.746   -3724.406
2673.979
engineloation_rear    1.159e+04   1638.588    7.071   0.000   8345.074
1.48e+04
enginetype_l          6988.6523   2392.209    2.921   0.004   2255.608
1.17e+04
enginetype_ohcf       5007.1933    847.950    5.905   0.000   3329.504
6684.883
enginetype_rotor      7418.4990   1461.610    5.076   0.000   4526.667
1.03e+04
fuelsystem_idi        1.294e+04   2869.216    4.512   0.000   7267.961
1.86e+04
car_company_bmw       8307.1217   1020.468    8.141   0.000   6288.100
1.03e+04
car_company_mazda     -1897.4175    825.214   -2.299   0.023   -3530.123
-264.712
car_company_mitsubishi -3077.6988    875.141   -3.517   0.001   -4809.186
-1346.211
car_company_peugeot   -1.097e+04   2629.810   -4.171   0.000   -1.62e+04
-5764.481
car_company_renault    -5273.2963   1661.724   -3.173   0.002   -8561.058
-1985.535
car_company_subaru     -6579.8676    951.545   -6.915   0.000   -8462.522
-4697.213
=====
Omnibus:                8.408   Durbin-Watson:                1.997
Prob(Omnibus):          0.015   Jarque-Bera (JB):            9.825
Skew:                   0.399   Prob(JB):                     0.00735
Kurtosis:               4.005   Cond. No.                     3.60e+16
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 2.23e-31. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

"""

```
[50]: vif=pd.DataFrame()
```

```
[51]: vif['Features'] = X_train_rfe1.columns
vif
```

```
[51]:
      Features
0      const
1      carwidth
2      enginesize
3      compressionratio
4      fueltype_gas
5      enginelocation_rear
6      enginetype_l
7      enginetype_ohcf
8      enginetype_rotor
9      fuelsystem_idi
10     car_company_bmw
11     car_company_mazda
12     car_company_mitsubishi
13     car_company_peugeot
14     car_company_renault
15     car_company_subaru
```

```
[52]: from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
[53]: vif['VIF'] = [variance_inflation_factor(X_train_rfe1.values, i) for i in
    ↪range(X_train_rfe1.shape[1])]

```

```
[54]: vif['VIF'] = round(vif['VIF'], 2)
```

```
[55]: vif.sort_values(by='VIF', ascending=False)
```

```
[55]:
      Features      VIF
4      fueltype_gas    inf
5      enginelocation_rear    inf
7      enginetype_ohcf    inf
9      fuelsystem_idi    inf
15     car_company_subaru    inf
3      compressionratio  42.32
13     car_company_peugeot  9.73
6      enginetype_l      8.99
1      carwidth          3.38
2      enginesize        3.30
8      enginetype_rotor  1.55
11     car_company_mazda  1.50
12     car_company_mitsubishi  1.20
10     car_company_bmw    1.12
14     car_company_renault  1.01
0      const            0.00
```

```
[56]: lm = LinearRegression()
      rfe2 = RFE(lm, 10)
```

```
rfe2.fit(X_train, y_train)
```

```
[56]: RFE(estimator=LinearRegression(), n_features_to_select=10)
```

```
[57]: rfe2.ranking_
```

```
[57]: array([46, 44, 33, 36,  1, 31, 22, 37,  1, 23, 38,  4, 28, 29, 42, 35,  3,
        21, 19, 17, 18, 20, 40, 45,  1, 27,  1, 25,  1, 47,  1, 41, 15,  2,
        48, 32, 24, 49, 14,  1, 39, 13, 11,  9, 30, 34,  6, 50,  5,  7,  1,
        12, 16,  1, 26,  1,  8, 10, 43])
```

```
[58]: col2 = X_train.columns[rfe2.support_]
X_train_rfe2 = X_train[col2]
X_train_rfe2 = sm.add_constant(X_train_rfe2)
lm2 = sm.OLS(y_train, X_train_rfe2).fit()
print(lm2.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          price    R-squared:                0.907
Model:                  OLS      Adj. R-squared:            0.901
Method:                 Least Squares    F-statistic:          144.3
Date:                  Sat, 12 Mar 2022    Prob (F-statistic):      3.98e-64
Time:                  20:09:02    Log-Likelihood:         -1314.2
No. Observations:      143    AIC:                     2648.
Df Residuals:          133    BIC:                     2678.
Df Model:               9
Covariance Type:       nonrobust
=====
```

```
=====
               coef      std err          t      P>|t|      [0.025
0.975]
-----
const          1.266e+04    235.167     53.845     0.000     1.22e+04
1.31e+04
carwidth       3586.8272    365.982      9.801     0.000     2862.929
4310.725
engine size    3739.2569    369.354     10.124     0.000     3008.690
4469.824
engine location_rear  1.131e+04    1735.038      6.519     0.000     7878.367
1.47e+04
engine type_l   7351.4558    2533.602      2.902     0.004     2340.089
1.24e+04
engine type_ohcf  5097.7417     897.846      5.678     0.000     3321.837
6873.646
engine type_rotor  5388.9829    1337.401      4.029     0.000     2743.655
8034.311
=====
```

car_company_bmw	8749.4458	1071.995	8.162	0.000	6629.081
1.09e+04					
car_company_peugeot	-9788.1466	2757.167	-3.550	0.001	-1.52e+04
-4334.578					
car_company_renault	-4866.7944	1757.035	-2.770	0.006	-8342.141
-1391.448					
car_company_subaru	-6212.4633	1003.232	-6.192	0.000	-8196.816
-4228.110					

```
=====
Omnibus:                    5.615    Durbin-Watson:                1.942
Prob(Omnibus):              0.060    Jarque-Bera (JB):          5.456
Skew:                      0.349    Prob(JB):                 0.0654
Kurtosis:                  3.655    Cond. No.                 2.00e+16
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 6.3e-31. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
[59]: vif=pd.DataFrame()
vif['Features'] = X_train_rfe2.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe2.values, i) for i in
↳range(X_train_rfe2.shape[1])]
vif.sort_values(by='VIF', ascending=False)
# showing vif
```

```
[59]:
```

	Features	VIF
3	enginolocation_rear	inf
5	enginetype_ohcf	inf
10	car_company_subaru	inf
8	car_company_peugeot	9.494513
4	enginetype_l	8.952555
2	enginesize	3.226114
1	carwidth	3.167483
0	const	1.307825
6	enginetype_rotor	1.150062
7	car_company_bmw	1.092396
9	car_company_renault	1.006775

```
[60]: X_train_rfe2.drop('car_company_subaru', axis=1, inplace=True)
```

```
[61]: X_train_rfe2 = sm.add_constant(X_train_rfe2)
lm2 = sm.OLS(y_train, X_train_rfe2).fit()
print(lm2.summary())
```

OLS Regression Results

```

=====
Dep. Variable:          price      R-squared:          0.907
Model:                  OLS        Adj. R-squared:       0.901
Method:                 Least Squares  F-statistic:        144.3
Date:                   Sat, 12 Mar 2022  Prob (F-statistic):    3.98e-64
Time:                   20:09:03    Log-Likelihood:     -1314.2
No. Observations:      143         AIC:                2648.
Df Residuals:          133         BIC:                2678.
Df Model:               9
Covariance Type:       nonrobust
=====

=====
                                coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const                1.266e+04    235.167      53.845      0.000    1.22e+04
1.31e+04
carwidth             3586.8272    365.982       9.801      0.000    2862.929
4310.725
engine              3739.2569    369.354     10.124      0.000    3008.690
4469.824
engineloation_rear  1.752e+04    2688.407       6.518      0.000    1.22e+04
2.28e+04
enginetype_l         7351.4558    2533.602       2.902      0.004    2340.089
1.24e+04
enginetype_ohcf     -1114.7216    784.120      -1.422      0.157   -2665.681
436.238
enginetype_rotor     5388.9829    1337.401       4.029      0.000    2743.655
8034.311
car_company_bmw       8749.4458    1071.995       8.162      0.000    6629.081
1.09e+04
car_company_peugeot -9788.1466    2757.167      -3.550      0.001   -1.52e+04
-4334.578
car_company_renault -4866.7944    1757.035      -2.770      0.006   -8342.141
-1391.448
=====
Omnibus:              5.615    Durbin-Watson:      1.942
Prob(Omnibus):        0.060    Jarque-Bera (JB):    5.456
Skew:                 0.349    Prob(JB):            0.0654
Kurtosis:             3.655    Cond. No.            23.8
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[62]: vif=pd.DataFrame()
vif['Features'] = X_train_rfe2.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe2.values, i) for i in
↳range(X_train_rfe2.shape[1])]
vif.sort_values(by='VIF', ascending=False)
# showing vif
```

```
[62]:
```

	Features	VIF
8	car_company_peugeot	9.494513
4	enginetype_l	8.952555
2	enginesize	3.226114
1	carwidth	3.167483
0	const	1.307825
3	enginolocation_rear	1.186865
6	enginetype_rotor	1.150062
5	enginetype_ohcf	1.117740
7	car_company_bmw	1.092396
9	car_company_renault	1.006775

```
[63]: X_train_rfe2.drop('enginetype_ohcf', axis=1, inplace=True)
```

```
[64]: X_train_rfe2 = sm.add_constant(X_train_rfe2)
lm2 = sm.OLS(y_train, X_train_rfe2).fit()
print(lm2.summary())
```

```

OLS Regression Results
=====
Dep. Variable:          price    R-squared:                0.906
Model:                  OLS      Adj. R-squared:            0.900
Method:                 Least Squares    F-statistic:            160.8
Date:                  Sat, 12 Mar 2022    Prob (F-statistic):      8.22e-65
Time:                  20:09:03    Log-Likelihood:          -1315.3
No. Observations:      143    AIC:                    2649.
Df Residuals:          134    BIC:                    2675.
Df Model:               8
Covariance Type:       nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025
0.975]					
const	1.256e+04	225.167	55.790	0.000	1.21e+04
1.3e+04					
carwidth	3575.4156	367.285	9.735	0.000	2848.989
4301.842					
enginesize	3788.6562	369.114	10.264	0.000	3058.614
4518.699					

engineloation_rear	1.642e+04	2583.971	6.355	0.000	1.13e+04
2.15e+04					
enginetype_l	7500.5753	2541.056	2.952	0.004	2474.810
1.25e+04					
enginetype_rotor	5552.1175	1337.536	4.151	0.000	2906.704
8197.531					
car_company_bmw	8800.1077	1075.476	8.183	0.000	6673.003
1.09e+04					
car_company_peugeot	-9839.1754	2767.416	-3.555	0.001	-1.53e+04
-4365.708					
car_company_renault	-4771.2505	1762.425	-2.707	0.008	-8257.020
-1285.481					

Omnibus:	5.533	Durbin-Watson:	1.944
Prob(Omnibus):	0.063	Jarque-Bera (JB):	5.168
Skew:	0.374	Prob(JB):	0.0755
Kurtosis:	3.555	Cond. No.	23.8

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[65]: vif=pd.DataFrame()
vif['Features'] = X_train_rfe2.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe2.values, i) for i in
↳range(X_train_rfe2.shape[1])]
vif.sort_values(by='VIF', ascending=False)
# showing vif
```

```
[65]:
```

	Features	VIF
7	car_company_peugeot	9.492904
4	enginetype_l	8.937210
2	enginesize	3.197559
1	carwidth	3.165959
0	const	1.189897
5	enginetype_rotor	1.141595
6	car_company_bmw	1.091189
3	engineloation_rear	1.088153
8	car_company_renault	1.005302

```
[66]: X_train_rfe2.drop('car_company_peugeot', axis=1, inplace=True)
```

```
[67]: X_train_rfe2 = sm.add_constant(X_train_rfe2)
lm2 = sm.OLS(y_train, X_train_rfe2).fit()
print(lm2.summary())
```

OLS Regression Results


```

=====
Dep. Variable:          price    R-squared:          0.897
Model:                  OLS      Adj. R-squared:       0.891
Method:                 Least Squares  F-statistic:        167.5
Date:                  Sat, 12 Mar 2022  Prob (F-statistic):    2.49e-63
Time:                  20:09:04   Log-Likelihood:     -1321.7
No. Observations:      143      AIC:                2659.
Df Residuals:          135      BIC:                2683.
Df Model:               7
Covariance Type:       nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const                1.254e+04    234.591     53.458     0.000    1.21e+04
1.3e+04
carwidth             3201.2839    366.745      8.729     0.000    2475.975
3926.593
engineize            3968.7863    381.057     10.415     0.000    3215.172
4722.400
engineloation_rear   1.599e+04    2690.180      5.946     0.000    1.07e+04
2.13e+04
enginetype_l         -963.1972    926.277     -1.040     0.300   -2795.089
868.694
enginetype_rotor     5781.1882    1392.391      4.152     0.000    3027.467
8534.910
car_company_bmw      8767.2497    1120.844      7.822     0.000    6550.566
1.1e+04
car_company_renault -4660.5447    1836.552     -2.538     0.012   -8292.678
-1028.411
=====

```

```

Omnibus:              10.615    Durbin-Watson:       1.972
Prob(Omnibus):        0.005    Jarque-Bera (JB):     11.115
Skew:                 0.570    Prob(JB):             0.00386
Kurtosis:             3.752    Cond. No.              16.6
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

[68]: vif=pd.DataFrame()
vif['Features'] = X_train_rfe2.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe2.values, i) for i in
↳range(X_train_rfe2.shape[1])]

```

```
vif.sort_values(by='VIF', ascending=False)
# showing vif
```

```
[68]:
```

	Features	VIF
2	enginesize	3.137317
1	carwidth	2.906076
0	const	1.189050
5	enginetype_rotor	1.138946
4	enginetype_l	1.093290
6	car_company_bmw	1.091108
3	enginelocation_rear	1.085818
7	car_company_renault	1.004989

```
[69]: X_train_rfe2.drop('enginetype_l', axis=1, inplace=True)
```

```
[70]: X_train_rfe2 = sm.add_constant(X_train_rfe2)
lm2 = sm.OLS(y_train, X_train_rfe2).fit()
print(lm2.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:                  0.896
Model:                            OLS    Adj. R-squared:              0.891
Method:                 Least Squares    F-statistic:                  195.2
Date:                  Sat, 12 Mar 2022    Prob (F-statistic):          2.92e-64
Time:                  20:09:04    Log-Likelihood:              -1322.3
No. Observations:          143    AIC:                          2659.
Df Residuals:              136    BIC:                          2679.
Df Model:                   6
Covariance Type:            nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const                1.247e+04    225.791     55.247     0.000     1.2e+04
1.29e+04
carwidth              3094.8228    352.270      8.785     0.000    2398.186
3791.459
enginesize            4048.8588    373.307     10.846     0.000    3310.621
4787.097
enginelocation_rear  1.589e+04    2688.922      5.908     0.000    1.06e+04
2.12e+04
enginetype_rotor      5943.7489    1384.001      4.295     0.000    3206.803
8680.695
car_company_bmw       8786.7032    1121.023      7.838     0.000    6569.813
1.1e+04

```

```
car_company_renault -4573.6566    1835.198    -2.492    0.014    -8202.872
-944.441
=====
Omnibus:                7.920    Durbin-Watson:                1.970
Prob(Omnibus):          0.019    Jarque-Bera (JB):          7.687
Skew:                   0.497    Prob(JB):                  0.0214
Kurtosis:               3.549    Cond. No.                  16.6
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[71]: vif=pd.DataFrame()
vif['Features'] = X_train_rfe2.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe2.values, i) for i in
↳range(X_train_rfe2.shape[1])]
vif.sort_values(by='VIF', ascending=False)
# showing vif
```

```
[71]:
```

	Features	VIF
2	enginesize	3.009203
1	carwidth	2.679605
4	enginetype_rotor	1.124589
0	const	1.100862
5	car_company_bmw	1.090804
3	enginelocation_rear	1.084154
6	car_company_renault	1.002908

7 Prediction

```
[72]: df_test[varlist]=sc.transform(df_test[varlist])
```

```
[73]: y_test = df_test.pop('price')
```

```
[78]: X_test = df_test
```

```
[79]: col2
```

```
[79]: Index(['carwidth', 'enginesize', 'enginelocation_rear', 'enginetype_l',
'enginetype_ohcf', 'enginetype_rotor', 'car_company_bmw',
'car_company_peugeot', 'car_company_renault', 'car_company_subaru'],
dtype='object')
```


```
[80]: X_test_rfe2 = X_test[col2]
```

```
[81]: X_test_rfe2 = X_test_rfe2.drop(['enginetype_1', 'car_company_peugeot',
    ↪ 'enginetype_ohcf', 'car_company_subaru'], axis=1)
```

```
[82]: X_test_rfe2 = sm.add_constant(X_test_rfe2)
```

```
[83]: y_pred = lm2.predict(X_test_rfe2)
```

```
[86]: plt.figure(figsize=(20,10))
plt.scatter(y_test, y_pred)
plt.show()
```



```
[87]: from sklearn.metrics import r2_score
```

```
[88]: r2_score(y_test, y_pred)
```

```
[88]: 0.8997211435182687
```

```
[89]: col2
```

```
[89]: Index(['carwidth', 'enginesize', 'enginelocation_rear', 'enginetype_1',
    'enginetype_ohcf', 'enginetype_rotor', 'car_company_bmw',
    'car_company_peugeot', 'car_company_renault', 'car_company_subaru'],
    dtype='object')
```

```
[91]: col2 = col2.drop(['enginetype_1', 'car_company_peugeot', 'enginetype_ohcf',
    ↪ 'car_company_subaru'])
```

```
[95]: plt.figure(figsize=(20,10))
sns.heatmap(df[col2].corr(), annot=True)
```

```
plt.show()
```

